

# Biomass\_Regression\_All\_Data\_two\_vars

May 24, 2017

This script builds a regression model to predict biomass based on lidar metrics. Data points at all scan angles are used.

```
In [1]: %matplotlib inline
import pandas, statsmodels, numpy, math, scipy.stats
from matplotlib import pyplot as plt
#from __future__ import division
from numpy import log as ln
import warnings
warnings.filterwarnings('ignore')
```

Data Prep Notes: The clipped las files were imported into ArcMap and converted to point shapefiles. The point shapefiles were each joined to the plot discs to add columns for plot ID and biomass. The joined points were then exported to .csvs in the directory CSVs\_ClippedAndJoined. The order of fields should run: x, y, z(height), angle, plotID, biomass. Points from discs with no associated biomass were dropped.

```
In [2]: firstfile = r"D:\Users\jkelly\CSVs_ClippedAndJoined\first_all_clipped_and_joi
first = pandas.read_csv(firstfile, header=0)
lastfile = r"D:\Users\jkelly\CSVs_ClippedAndJoined\last_all_clipped_and_joi
last = pandas.read_csv(lastfile, header=0)
groundfile = r"D:\Users\jkelly\CSVs_ClippedAndJoined\Grounds_clipped_and_joi
ground = pandas.read_csv(groundfile, header=0)
```

```
In [3]: first.columns=['X1', 'Y1', 'X', 'Y', 'height', 'angle', 'PlotID', 'biomass']
del first['X1']
del first['Y1']
last.columns=['X1', 'Y1', 'X', 'Y', 'height', 'angle', 'PlotID', 'biomass']
del last['X1']
del last['Y1']
ground.columns=['X1', 'Y1', 'X', 'Y', 'height', 'angle', 'PlotID', 'biomass']
del ground['X1']
del ground['Y1']
```

```
In [4]: first=first.sort('PlotID')
last=last.sort('PlotID')
ground=ground.sort('PlotID')
```

The following function produces a series of metrics summarizing the lidar height measurements on each inventory plot for first, last and ground returns: mean, coefficient of variation, 10th, 50th and 90th deciles for height and density fraction, following Gobakken et al, 2012.

```
In [5]: def heightstats(SERIES, I):
    # function to compute "cloud metrics" on a height or elevation series
    totalechoes = len(SERIES)
    MEAN = numpy.mean(SERIES)
    CV = scipy.stats.variation(SERIES)
    #coeff of var (ratio of biased standard dev to mean)
    H10 = numpy.percentile(SERIES, 10)
    H10series = SERIES[SERIES > H10]
    H10count = len(H10series)
    D10 = H10count/totalechoes
    # "canopy density" = the proportion of laser
    # hits above fraction to total # of echoes
    # described in Gobakken et al. 2012, p. 447
    H50 = numpy.percentile(SERIES, 50)
    H50series = SERIES[SERIES > H50]
    H50count = len(H50series)
    D50 = H50count/totalechoes
    H90 = numpy.percentile(SERIES, 90)
    H90series = SERIES[SERIES > H90]
    H90count = len(H90series)
    D90 = H90count/totalechoes

    outlist = [I, MEAN, CV, H10, H50, H90, D10, D50, D90]

    return(outlist)
```

The next section builds a dataframe associating each plot with its biomass, elevation, and computed summary statistics. The plot-level biomass numbers given in the data are in units of Mg/ha, which are here converted to the actual mass on a the 0.04 ha plots.

```
In [6]: heightlist = []
    for i in first.PlotID.unique():
        # subset corresponding to ith plot id
        L = last[last.PlotID == i]
        F = first[first.PlotID == i]
        G = ground[ground.PlotID == float(i)]

        # just height column from subset
        Last_heightseries = L.height
        First_heightseries = F.height
        elevseries = G.height

        # just biomass column from subset
        biomass = L.biomass*0.04
        # the 0.04 is to convert Mg/ha to Mg/400m^2 plot
```

```

if len>Last_heightseries)>=1 and len(First_heightseries)>=1:
    # ^^ gets rid of empty height series

    # fetch "cloud metrics"
    Last_stats = heightstats>Last_heightseries, i)
    First_stats = heightstats(First_heightseries, i)

    if len(elevseries) > 0:
        Gmean = numpy.mean(elevseries)
    else: Gmean = -1

    # build a list for next row in dataframe
    stats = Last_stats
    for item in First_stats:
        stats.append(item)
    b = numpy.mean(biomass)
    stats.append(b)
    stats.append(Gmean)
    if b>0: heightlist.append(stats) #gets rid of NaN

heightarray = numpy.array(heightlist) # build the dataframe
df = pandas.DataFrame(heightarray)
columnlist = ['PlotID', 'L_mean', 'L_cv', 'L_h10', 'L_h50',
              'L_h90', 'L_d10', 'L_d50', 'L_d90', 'PlotID_again',
              'F_mean', 'F_cv', 'F_h10', 'F_h50', 'F_h90', 'F_d10',
              'F_d50', 'F_d90', 'Biomass', 'elev']

df.columns = columnlist
# delete duplicate columns
del df['PlotID_again']

```

In [7]: df.describe()

```

Out[7]:

```

	PlotID	L_mean	L_cv	L_h10	L_h50	L_h90
count	154.000000	154.000000	154.000000	154.000000	154.000000	154.000000
mean	93.116883	2.882272	1.508728	0.002792	0.931591	8.772792
std	50.724934	1.713424	0.433447	0.012153	1.381876	4.358676
min	2.000000	0.206482	0.666363	0.000000	0.040000	0.240000
25%	55.250000	1.588867	1.202409	0.000000	0.140000	5.492500
50%	96.000000	2.550631	1.426635	0.000000	0.270000	8.385000
75%	135.750000	4.170365	1.715680	0.000000	0.993750	11.716500
max	178.000000	7.915523	2.857065	0.120000	6.380000	19.804000

	L_d10	L_d50	L_d90	F_mean	F_cv	F_h10
count	154.000000	154.000000	154.000000	154.000000	154.000000	154.000000
mean	0.842814	0.496161	0.100046	11.446377	0.504331	3.730400

std	0.039673	0.005197	0.000534	3.564578	0.289637	4.2103
min	0.729089	0.468468	0.098137	1.220464	0.166920	0.0000
25%	0.820128	0.494567	0.099872	9.087133	0.348159	0.1017
50%	0.847703	0.498274	0.100080	10.943383	0.431127	1.9790
75%	0.873282	0.499613	0.100323	13.433398	0.584139	6.4945
max	0.899889	0.500000	0.102804	21.774059	2.199853	16.2100

	F_h50	F_h90	F_d10	F_d50	F_d90	Bioma
count	154.000000	154.000000	154.000000	154.000000	154.000000	154.0000
mean	12.367338	17.210695	0.896085	0.499204	0.099676	5.7529
std	4.041083	3.812580	0.013928	0.002669	0.000431	2.8232
min	0.070000	4.961000	0.789216	0.468215	0.097949	0.9120
25%	9.820000	14.464000	0.898350	0.499200	0.099473	3.6420
50%	12.105000	16.930000	0.899754	0.499585	0.099782	5.1960
75%	15.020000	19.537500	0.899880	0.499881	0.100000	6.9258
max	23.000000	26.640000	0.900173	0.500000	0.100654	17.4634

	elev
count	154.000000
mean	3.118651
std	1.874205
min	0.194529
25%	1.573210
50%	2.635528
75%	4.435760
max	8.226820

Then a regression model is developed and its statistics are displayed. Though Gobakken et al developed different models for different forest classes, in each case both the response and explanatory variables were log-transformed. However, an exhaustive trial of all one-, two- and three-variable combinations revealed the "best" two-variable model –in terms of  $R^2$ , as well as AIC compared on identical variables – to be one which log-transforms the 50th percentile cutoff for height of first returns and also includes the untransformed value for first-return mean. To rule out (multi) collinearity, models with a condition number greater than 50 were excluded from consideration.

```
In [19]: def formula_tester(VAR1, VAR2):
        try:
            #three variables three logs (not a contender)
            #Formula = "ln(Y) ~ 1 + ln(" + VAR1 + ") + ln(" + VAR2 + ") + ln(" + VAR3 + ")"

            # three variables two logs - best model so far
            #Formula = "ln(Y) ~ 1 + ln(" + VAR1 + ") + ln(" + VAR2 + ") + " + VAR3

            # three variables one log -- good
            #Formula = "ln(Y) ~ 1 + ln(" + VAR1 + ") + " + VAR2 + " + " + VAR3

            # three variables no log
            #Formula = "ln(Y) ~ 1 + " + VAR1 + " + " + VAR2 + " + " + VAR3
```

```

# one variable log (not a contender)
#Formula = "ln(Y) ~ 1 + ln( " + VAR1 + ")"

# two variables two logs (not a contender)
#Formula = "ln(Y) ~ 1 + ln( " + VAR1 + ") + " + "ln( " + VAR2 + ")"

#two variables one log -- not bad
Formula = "ln(Y) ~ 1 + ln(" + VAR1 + ") + " + VAR2

result = sm.ols(formula=Formula, data=df).fit()
if result.condition_number < 50 and result.rsquared>0.7:
    return(result.rsquared, result.aic, VAR1, VAR2)
    #return(result.rsquared, result.aic, VAR1, VAR2) #two variables
else:
    return()
except:
    return()

```

In [20]: `import statsmodels.formula.api as sm`

```

Y = df['Biomass']
variable_list = ["df['F_mean']", "df['L_mean']", "df['F_cv']", "df['L_cv']",
                 "df['F_h10']", "df['L_h10']", "df['F_h50']", "df['L_h50']",
                 "df['F_h90']", "df['L_h90']", "df['F_d10']", "df['L_d10']",
                 "df['F_d50']", "df['L_d50']", "df['F_d90']", "df['L_d90']",
                 "df['elev']"]

regression_parameters_list = []
for variable1 in variable_list:
    for variable2 in variable_list:

        regression_parameters_list.append(formula_tester(variable1, variable2))
        parameters = formula_tester(variable1, variable2)

        #remove empty tuples and report best r-squareds
        if len(parameters) > 1 and parameters[0]>0.7:
            regression_parameters_list.append(parameters)
sorted(filter(lambda a: a != (), regression_parameters_list))
#regression_parameters_list

```

Out[20]: [(0.7136724104353942, 25.914568617848545, "df['L\_h90']", "df['F\_mean']"),  
(0.71536310143767801, 25.002542806256713, "df['F\_cv']", "df['F\_mean']"),  
(0.71811908422823867, 23.504179693283731, "df['elev']", "df['F\_mean']"),  
(0.71849397496043177, 23.299229322244571, "df['L\_mean']", "df['F\_mean']"),  
(0.76808072198590027, -6.5405004686309667, "df['F\_h50']", "df['F\_mean']"),  
(0.77657977233883646, -12.290081906506089, "df['F\_h50']", "df['F\_h50'])]

This was the second best regression model with 2 independent variables, condition number < 50.

The best model produced unacceptable overestimates on non-forested areas.

```
In [21]: var1 = df['F_h50']
        var2 = df['F_mean']

        var1str = "ln(df['F_h50'])"
        var2str = "df['F_mean']"

        Formula = "ln(Y) ~ 1 + " + var1str + " + " + var2str
        result = sm.ols(formula=Formula, data=df).fit()
        print(result.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  ln(Y)      R-squared:                0.768
Model:                        OLS          Adj. R-squared:           0.765
Method:                    Least Squares   F-statistic:              250.0
Date:                Wed, 24 May 2017      Prob (F-statistic):       1.21e-48
Time:                  10:45:24            Log-Likelihood:           6.2703
No. Observations:          154             AIC:                    -6.541
Df Residuals:              151             BIC:                    2.570
Df Model:                   2
Covariance Type:            nonrobust
=====

```

	coef	std err	t	P> t	[95.0% Conf. Int.]
Intercept	0.4665	0.068	6.857	0.000	0.332 0.601
ln(df['F_h50'])	-0.2049	0.034	-5.998	0.000	-0.272 -0.137
df['F_mean']	0.1449	0.007	19.755	0.000	0.130 0.159

```

=====
Omnibus:                    39.687      Durbin-Watson:           1.609
Prob(Omnibus):              0.000      Jarque-Bera (JB):        206.187
Skew:                      -0.761      Prob(JB):                1.69e-45
Kurtosis:                   8.460      Cond. No.:               45.0
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [22]: result.bse
```

```
Out[22]: Intercept          0.068035
        ln(df['F_h50'])      0.034156
        df['F_mean']         0.007336
        dtype: float64
```

```

In [25]: intercept = result.params[0]
        b1 = result.params[1]
        b2 = result.params[2]

        intercept, b1, b2

Out[25]: (0.46653258757917682, -0.20488443892984287, 0.14492814972141516)

In [27]: f = open('parameters_all_two.csv', 'w')
        with f:
            f.write(str(intercept)+' \n ')
            f.write(str(b1)+ ' \n ')
            f.write(str(b2)+ '\n' )

        f.close()


```

Inspecting the data and residuals:

```

In [29]: R = result.resid
        plt.plot(Y, R, '.')
        plt.hlines(0,0, 20,colors=u'r')
        plt.show()

```



```

In [30]: df.mean()

Out[30]: PlotID      93.116883
        L_mean      2.882272

```

```

L_cv      1.508728
L_h10     0.002792
L_h50     0.931591
L_h90     8.772786
L_d10     0.842814
L_d50     0.496161
L_d90     0.100046
F_mean    11.446377
F_cv      0.504331
F_h10     3.730494
F_h50     12.367338
F_h90     17.210695
F_d10     0.896085
F_d50     0.499204
F_d90     0.099676
Biomass   5.752972
elev      3.118651
dtype: float64

```

```
In [31]: #df.to_csv("ALL_plots_yhat")
```

```
In [33]: df.describe()
```

```

Out[33]:

```

	PlotID	L_mean	L_cv	L_h10	L_h50	L_h90
count	154.000000	154.000000	154.000000	154.000000	154.000000	154.000000
mean	93.116883	2.882272	1.508728	0.002792	0.931591	8.772786
std	50.724934	1.713424	0.433447	0.012153	1.381876	4.358867
min	2.000000	0.206482	0.666363	0.000000	0.040000	0.240000
25%	55.250000	1.588867	1.202409	0.000000	0.140000	5.492000
50%	96.000000	2.550631	1.426635	0.000000	0.270000	8.385000
75%	135.750000	4.170365	1.715680	0.000000	0.993750	11.716000
max	178.000000	7.915523	2.857065	0.120000	6.380000	19.804000

	L_d10	L_d50	L_d90	F_mean	F_cv	F_h10
count	154.000000	154.000000	154.000000	154.000000	154.000000	154.000000
mean	0.842814	0.496161	0.100046	11.446377	0.504331	3.730494
std	0.039673	0.005197	0.000534	3.564578	0.289637	4.210000
min	0.729089	0.468468	0.098137	1.220464	0.166920	0.000000
25%	0.820128	0.494567	0.099872	9.087133	0.348159	0.101000
50%	0.847703	0.498274	0.100080	10.943383	0.431127	1.979000
75%	0.873282	0.499613	0.100323	13.433398	0.584139	6.494000
max	0.899889	0.500000	0.102804	21.774059	2.199853	16.210000

	F_h50	F_h90	F_d10	F_d50	F_d90	Biomass
count	154.000000	154.000000	154.000000	154.000000	154.000000	154.000000
mean	12.367338	17.210695	0.896085	0.499204	0.099676	5.752972
std	4.041083	3.812580	0.013928	0.002669	0.000431	2.823000
min	0.070000	4.961000	0.789216	0.468215	0.097949	0.912000



25%	9.820000	14.464000	0.898350	0.499200	0.099473	3.642
50%	12.105000	16.930000	0.899754	0.499585	0.099782	5.190
75%	15.020000	19.537500	0.899880	0.499881	0.100000	6.925
max	23.000000	26.640000	0.900173	0.500000	0.100654	17.463

	elev
count	154.000000
mean	3.118651
std	1.874205
min	0.194529
25%	1.573210
50%	2.635528
75%	4.435760
max	8.226820

In [ ]: