

# Biomass\_Regression\_All\_Data\_three\_vars

May 24, 2017

This script builds a regression model to predict biomass based on lidar metrics. Data points at all scan angles are used.

```
In [1]: %matplotlib inline
import pandas, statsmodels, numpy, math, scipy.stats
from matplotlib import pyplot as plt
#from __future__ import division
from numpy import log as ln
import warnings
warnings.filterwarnings('ignore')
```

Data Prep Notes: The clipped las files were imported into ArcMap and converted to point shapefiles. The point shapefiles were each joined to the plot discs to add columns for plot ID and biomass. The joined points were then exported to .csvs in the directory CSVs\_ClippedAndJoined. The order of fields should run: x, y, z(height), angle, plotID, biomass. Points from discs with no associated biomass were dropped.

```
In [2]: firstfile = r"D:\Users\jkelly\CSVs_ClippedAndJoined\first_all_clipped_and_joi
first = pandas.read_csv(firstfile, header=0)
lastfile = r"D:\Users\jkelly\CSVs_ClippedAndJoined\last_all_clipped_and_joi
last = pandas.read_csv(lastfile, header=0)
groundfile = r"D:\Users\jkelly\CSVs_ClippedAndJoined\Grounds_clipped_and_joi
ground = pandas.read_csv(groundfile, header=0)
```

```
In [3]: first.columns=['X1', 'Y1', 'X', 'Y', 'height', 'angle', 'PlotID', 'biomass']
del first['X1']
del first['Y1']
last.columns=['X1', 'Y1', 'X', 'Y', 'height', 'angle', 'PlotID', 'biomass']
del last['X1']
del last['Y1']
ground.columns=['X1', 'Y1', 'X', 'Y', 'height', 'angle', 'PlotID', 'biomass']
del ground['X1']
del ground['Y1']
```

```
In [4]: first=first.sort('PlotID')
last=last.sort('PlotID')
ground=ground.sort('PlotID')
```

The following function produces a series of metrics summarizing the lidar height measurements on each inventory plot for first, last and ground returns: mean, coefficient of variation, 10th, 50th and 90th deciles for height and density fraction, following Gobakken et al, 2012.

```
In [5]: def heightstats(SERIES, I):
    # function to compute "cloud metrics" on a height or elevation series
    totalechoes = len(SERIES)
    MEAN = numpy.mean(SERIES)
    CV = scipy.stats.variation(SERIES)
    #coeff of var (ratio of biased standard dev to mean)
    H10 = numpy.percentile(SERIES, 10)
    H10series = SERIES[SERIES > H10]
    H10count = len(H10series)
    D10 = H10count/totalechoes
    # "canopy density" = the proportion of laser
    # hits above fraction to total # of echoes
    # described in Gobakken et al. 2012, p. 447
    H50 = numpy.percentile(SERIES, 50)
    H50series = SERIES[SERIES > H50]
    H50count = len(H50series)
    D50 = H50count/totalechoes
    H90 = numpy.percentile(SERIES, 90)
    H90series = SERIES[SERIES > H90]
    H90count = len(H90series)
    D90 = H90count/totalechoes

    outlist = [I, MEAN, CV, H10, H50, H90, D10, D50, D90]

    return(outlist)
```

The next section builds a dataframe associating each plot with its biomass, elevation, and computed summary statistics. The plot-level biomass numbers given in the data are in units of Mg/ha, which are here converted to the actual mass on a the 0.04 ha plots.

```
In [6]: heightlist = []
    for i in first.PlotID.unique():
        # subset corresponding to ith plot id
        L = last[last.PlotID == i]
        F = first[first.PlotID == i]
        G = ground[ground.PlotID == float(i)]

        # just height column from subset
        Last_heightseries = L.height
        First_heightseries = F.height
        elevseries = G.height

        # just biomass column from subset
        biomass = L.biomass*0.04
        # the 0.04 is to convert Mg/ha to Mg/400m^2 plot
```

```

if len>Last_heightseries)>=1 and len(First_heightseries)>=1:
    # ^^ gets rid of empty height series

    # fetch "cloud metrics"
    Last_stats = heightstats>Last_heightseries, i)
    First_stats = heightstats(First_heightseries, i)

    if len(elevseries) > 0:
        Gmean = numpy.mean(elevseries)
    else: Gmean = -1

    # build a list for next row in dataframe
    stats = Last_stats
    for item in First_stats:
        stats.append(item)
    b = numpy.mean(biomass)
    stats.append(b)
    stats.append(Gmean)
    if b>0: heightlist.append(stats) #gets rid of NaN

heightarray = numpy.array(heightlist) # build the dataframe
df = pandas.DataFrame(heightarray)
columnlist = ['PlotID', 'L_mean', 'L_cv', 'L_h10', 'L_h50',
              'L_h90', 'L_d10', 'L_d50', 'L_d90', 'PlotID_again',
              'F_mean', 'F_cv', 'F_h10', 'F_h50', 'F_h90', 'F_d10',
              'F_d50', 'F_d90', 'Biomass', 'elev']

df.columns = columnlist
# delete duplicate columns
del df['PlotID_again']

```

In [19]: df.describe()

```

Out[19]:

```

	PlotID	L_mean	L_cv	L_h10	L_h50	L_h90
count	154.000000	154.000000	154.000000	154.000000	154.000000	154.000000
mean	93.116883	2.882272	1.508728	0.002792	0.931591	8.772000
std	50.724934	1.713424	0.433447	0.012153	1.381876	4.358000
min	2.000000	0.206482	0.666363	0.000000	0.040000	0.240000
25%	55.250000	1.588867	1.202409	0.000000	0.140000	5.492000
50%	96.000000	2.550631	1.426635	0.000000	0.270000	8.385000
75%	135.750000	4.170365	1.715680	0.000000	0.993750	11.716000
max	178.000000	7.915523	2.857065	0.120000	6.380000	19.804000

	L_d10	L_d50	L_d90	F_mean	...	F_h10
count	154.000000	154.000000	154.000000	154.000000	...	154.000000
mean	0.842814	0.496161	0.100046	11.446377	...	3.730000

std	0.039673	0.005197	0.000534	3.564578	...	4.210
min	0.729089	0.468468	0.098137	1.220464	...	0.000
25%	0.820128	0.494567	0.099872	9.087133	...	0.101
50%	0.847703	0.498274	0.100080	10.943383	...	1.979
75%	0.873282	0.499613	0.100323	13.433398	...	6.494
max	0.899889	0.500000	0.102804	21.774059	...	16.210

	F_h50	F_h90	F_d10	F_d50	F_d90	Biom
count	154.000000	154.000000	154.000000	154.000000	154.000000	154.000
mean	12.367338	17.210695	0.896085	0.499204	0.099676	5.752
std	4.041083	3.812580	0.013928	0.002669	0.000431	2.823
min	0.070000	4.961000	0.789216	0.468215	0.097949	0.912
25%	9.820000	14.464000	0.898350	0.499200	0.099473	3.642
50%	12.105000	16.930000	0.899754	0.499585	0.099782	5.190
75%	15.020000	19.537500	0.899880	0.499881	0.100000	6.925
max	23.000000	26.640000	0.900173	0.500000	0.100654	17.463

	elev	Yhat	resid
count	154.000000	154.000000	154.000000
mean	3.118651	5.673073	0.079899
std	1.874205	2.902675	1.235315
min	0.194529	2.434643	-3.777283
25%	1.573210	3.766871	-0.531332
50%	2.635528	4.686643	0.058783
75%	4.435760	6.479547	0.715873
max	8.226820	19.683735	4.903069

[8 rows x 21 columns]

Then a regression model is developed and its statistics are displayed. Though Gobakken et al developed different models for different forest classes, in each case both the response and explanatory variables were log-transformed. However, an exhaustive trial of all one-, two- and three-variable combinations revealed the "best" model –in terms of  $R^2$ , as well as AIC compared on identical variables – to be one which log-transforms the 50th percentile cutoff for height of first returns and the mean of last returns, but also includes the untransformed value for 50th percentile. To rule out (multi) collinearity, models with a condition number greater than 50 were excluded from consideration.

```
In [60]: def formula_tester(VAR1, VAR2, VAR3):
        try:
            #three variables three logs (not a contender)
            #Formula = "ln(Y) ~ 1 + ln(" + VAR1 + ") + ln(" + VAR2 + ") + ln(" + VAR3 + ")"

            # three variables two logs - best model so far
            Formula = "ln(Y) ~ 1 + ln(" + VAR1 + ") + ln(" + VAR2 + ") + " + VAR3

            # three variables one log -- good
            #Formula = "ln(Y) ~ 1 + ln(" + VAR1 + ") + " + VAR2 + " + " + VAR3
```

```

# three variables no log
#Formula = "ln(Y) ~ 1 + " + VAR1 + " + " + VAR2 + " + " + VAR3

# one variable log (not a contender)
#Formula = "ln(Y) ~ 1 + ln( " + VAR1 + " )"

# two variables two logs (not a contender)
#Formula = "ln(Y) ~ 1 + ln( " + VAR1 + " ) + " + "ln( " + VAR2 + " )"

#two variables one log -- not bad
#Formula = "ln(Y) ~ 1 + ln( " + VAR1 + " ) + " + VAR2

result = sm.ols(formula=Formula, data=df).fit()
if result.condition_number < 50 and result.rsquared>0.77:
    return(result.rsquared, result.aic, VAR1, VAR2, VAR3)
    #return(result.rsquared, result.aic, VAR1, VAR2) #two variables
else:
    return()
except:
    return()

```

```
In [61]: import statsmodels.formula.api as sm
```

```

Y = df['Biomass']
variable_list = ["df['F_mean']", "df['L_mean']", "df['F_cv']", "df['L_cv']",
                 "df['F_h10']", "df['L_h10']", "df['F_h50']", "df['L_h50']",
                 "df['F_h90']", "df['L_h90']", "df['F_d10']", "df['L_d10']",
                 "df['F_d50']", "df['L_d50']", "df['F_d90']", "df['L_d90']",
                 "df['elev']"]

regression_parameters_list = []
for variable1 in variable_list:
    for variable2 in variable_list:
        for variable3 in variable_list:
            regression_parameters_list.append(formula_tester(variable1, variable2, variable3))
            parameters = formula_tester(variable1, variable2, variable3)

#remove empty tuples and report best r-squareds
if len(parameters) > 1 and parameters[0]>0.7:
    regression_parameters_list.append(parameters)
sorted(filter(lambda a: a != (), regression_parameters_list))
#regression_parameters_list

```

```

Out[61]: [(0.77657977233883646,
          -12.290081906506089,
          "df['F_h50']",
          "df['F_h50']",
          "df['F_h50']"),

```

```
(0.78999759828180161,
-19.828142703340745,
"df['F_h50']",
"df['elev']",
"df['F_h50']"),
(0.78999759828180161,
-19.828142703340745,
"df['elev']",
"df['F_h50']",
"df['F_h50']"),
(0.79410704800067033,
-22.871580665461238,
"df['F_h50']",
"df['L_mean']",
"df['F_h50']"),
(0.79410704800067033,
-22.871580665461181,
"df['L_mean']",
"df['F_h50']",
"df['F_h50']")]
```

This was the best regression model with 3 independent variables, condition number < 50.

```
In [24]: var1 = df['L_mean']
var2 = df['F_h50']
var3 = df['F_h50']

var1str = "ln(df['L_mean'])"
var2str = "ln(df['F_h50'])"
var3str = "df['F_h50']"

Formula = "ln(Y) ~ 1 + " + var1str + " + " + var2str + " + " + var3str
result = sm.ols(formula=Formula, data=df).fit()
print(result.summary())
```

```

OLS Regression Results
=====
Dep. Variable:          ln(Y)      R-squared:                0.794
Model:                  OLS        Adj. R-squared:           0.790
Method:                 Least Squares   F-statistic:             192.8
Date:                   Tue, 16 May 2017   Prob (F-statistic):       2.92e-51
Time:                   11:38:07         Log-Likelihood:           15.436
No. Observations:       154             AIC:                     -22.87
Df Residuals:           150             BIC:                     -10.72
Df Model:                3
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[95.0% Conf. Int.]
-----					

Intercept	0.6916	0.064	10.811	0.000	0.565	0.81
ln(df['L_mean'])	0.1349	0.038	3.573	0.000	0.060	0.21
ln(df['F_h50'])	-0.3184	0.037	-8.532	0.000	-0.392	-0.24
df['F_h50']	0.1284	0.009	14.769	0.000	0.111	0.14

```
=====
Omnibus:                    54.093    Durbin-Watson:                1.804
Prob(Omnibus):              0.000    Jarque-Bera (JB):          271.744
Skew:                      -1.155    Prob(JB):                 9.81e-60
Kurtosis:                  9.084    Cond. No.                 49.3
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly sp

```
In [34]: result.bse
```

```
Out[34]: Intercept          0.063969
         ln(df['L_mean'])    0.037759
         ln(df['F_h50'])     0.037320
         df['F_h50']         0.008696
         dtype: float64
```


```
In [25]: intercept = result.params[0]
         b1 = result.params[1]
         b2 = result.params[2]
         b3 = result.params[3]
         intercept, b1, b2, b3
```

```
Out[25]: (0.69157584075617873,
         0.1349292756691241,
         -0.31840448145553918,
         0.12843720787213342)
```


```
In [64]: f = open('parameters_all.csv', 'w')
         with f:
             f.write(str(intercept)+' \n ')
             f.write(str(b1)+ ' \n ' )
             f.write(str(b2)+ '\n' )
             f.write(str(b3)+ '\n')
         f.close()
```

Inspecting the data and residuals:

```
In [27]: #X = ln(df.F_h50)
         X = df.F_mean
         Y = df.Biomass
         plt.plot(X, Y, '.')
         plt.show()
```



```
In [28]: R = result.resid
plt.plot(Y, R, '.')
plt.hlines(0,0, 20,colors=u'r')
plt.show()
```





```
In [30]: df.mean()
```

```
Out[30]: PlotID      93.116883  
         L_mean      2.882272  
         L_cv        1.508728  
         L_h10       0.002792  
         L_h50       0.931591  
         L_h90       8.772786  
         L_d10       0.842814  
         L_d50       0.496161  
         L_d90       0.100046  
         F_mean     11.446377  
         F_cv        0.504331  
         F_h10       3.730494  
         F_h50     12.367338  
         F_h90     17.210695  
         F_d10       0.896085  
         F_d50       0.499204  
         F_d90       0.099676  
         Biomass     5.752972  
         elev       3.118651  
         Yhat        0.112888  
         resid       5.640083  
         dtype: float64
```

```
In [18]: #df.to_csv("ALL_plots_yhat")
```